

Processes and Threads

CSCI 3431: Operating Systems

Agenda

- Assignment 1 due September 29!
 - Linux/Unix environment available on the CS machine
 - `ssh cs.smu.ca`
- Today's lecture
 - Thread Model
 - Java thread libraries
 - Issues related to multithreaded programming
- Reading: Sections 4.1-4.5
- Next :Java thread libraries, synchronization

Just checking....


- What is a process?
- What information does the OS need to manage a process?
- What data structures are used?
- What is the difference between I/O-bound and CPU-bound processes?
- Why do we care?

Motivation

- How do we handle complex applications?

Processes and Threads

- Single Threaded Approaches:
 - A single thread of execution per process, in which the concept of a thread is not recognized, is referred to as a single-threaded approach
 - MS-DOS is an example
- Multithreaded Approaches:
 - One process with multiple threads (JRE)
 - Multiple processes , each supporting multiple threads (Windows, Solaris, modern versions of UNIX)



Processes and Thread

Resource Ownership

- **Process image:** collection of program, data, stack, and attributes defined in the process control block.
- The OS performs a protection function to prevent unwanted interference between processes with respect to resources.

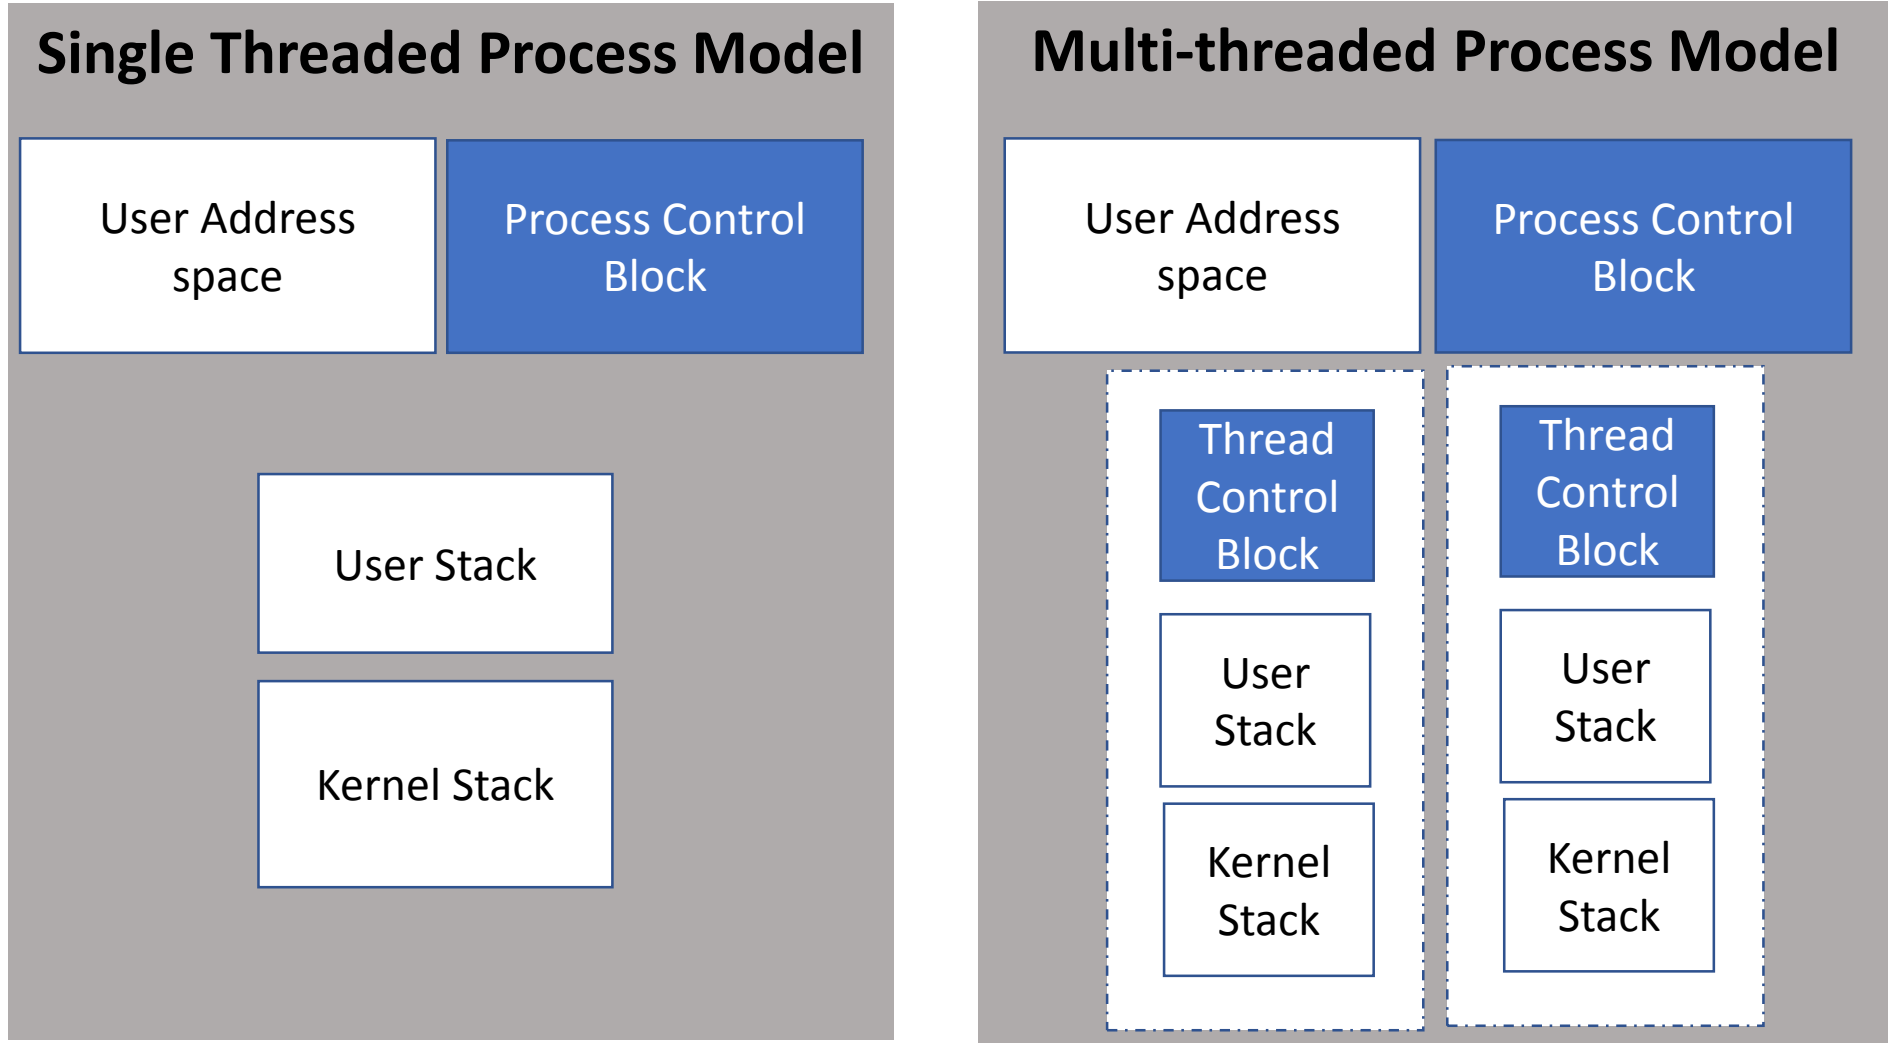
Scheduling/Execution

- Execution state (Running, Ready, etc.)
- The OS takes care of dispatching and scheduling a process
 - Priority
 - Execution may be interleaved with that of other processes

Processes and Threads

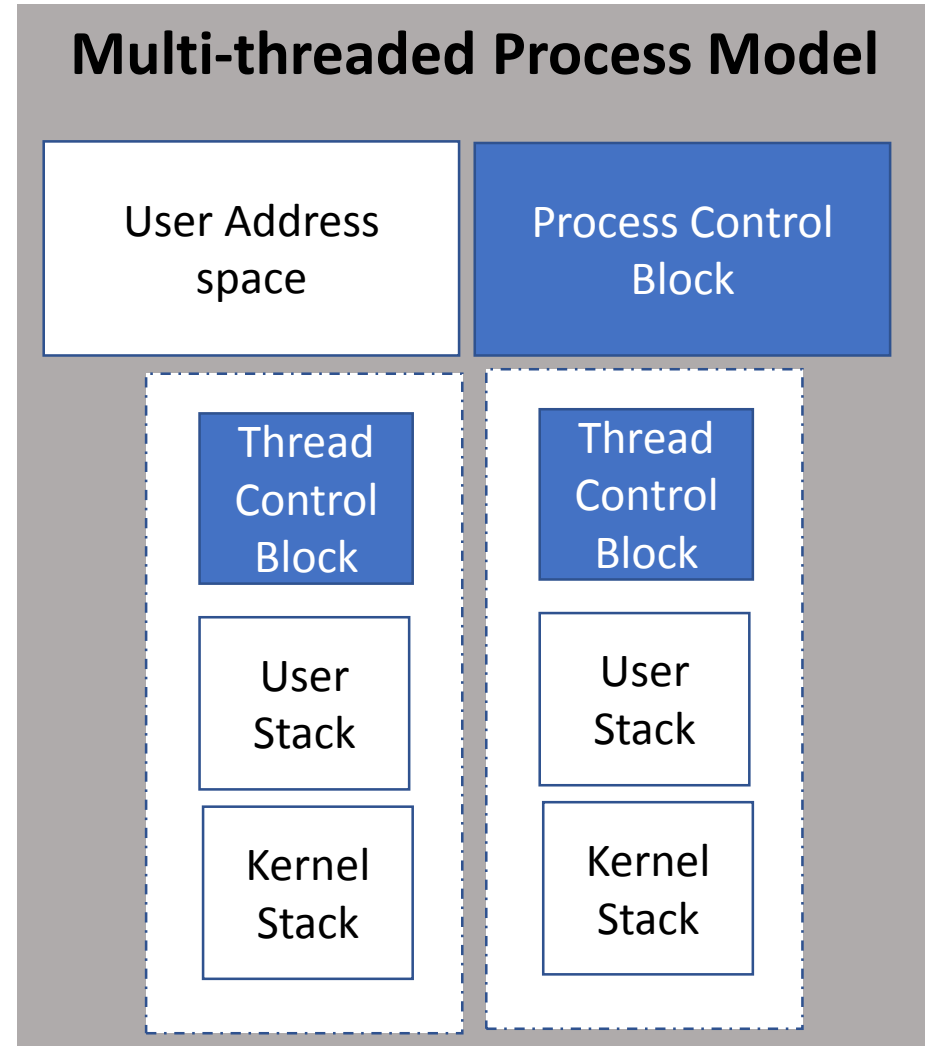
- **Process** or **task** refers to the unit of resource ownership
- **Thread** or **lightweight process** refers to the unit of dispatching
- **Multithreading** refers to the ability of an OS to support multiple, concurrent execution within a single process

Single and Multi-threaded Processes



One or More Threads in a Process

- Each thread has:
 - Access to the memory and resources of its process, shared with other threads in that process
 - CPU context
 - CPU registers
 - PC
 - Its own execution **stack**
 - Local variables
 - Thread Control Block (TCB)
 - Thread context (CPU context)
 - State (Running, Ready, etc.)



Key Benefits of Threads

- Responsiveness
- Resource Sharing
- Economy
- Scalability

Key Benefits of Threads

- Responsiveness
 - Loading webpage content made faster using threads
- Discuss the performance of Single Thread Web Browser with a colleague.
 - You click on a link
 - Decide to click on the stop button before you received a response

Key Benefits of Threads

- Resource Sharing
 - Process state and resources are shared among process threads
 - Same address space with access to same data
 - E.g., If one thread opens a file with read privileges, other threads in the same process can also read from that file
 - Enhanced communication between programs
- Discuss the implementation of a word processor
 - Handling Input
 - Formatting and displaying content
 - Backup

Key Benefits of Threads

- Economy
 - It takes less time to create/terminate a thread than a process (10x-100x faster)
 - The memory is already allocated for the process
 - No data or code needs to be copied
 - Thread creation is cheap
- It takes less time to switch between two threads within the same process than to switch between processes.

Key Benefits of Threads

- Scalability
 - Threads are already designed to run concurrently
 - Threads can be scheduled to run on multiple CPUs
 - The OS is responsible for scheduling threads
 - The application cannot tell if it is running on a single or multiprocessor system
 - Work can be distributed among multiple processors (when available)
- Increased parallelism on multiprocessor machines


Concurrent Processes

- Most OS provide process and thread abstractions to allow creation of multithread and multi-process application
- Hence, we talk about **concurrent processes**

Concurrent Processes


Concurrent Execution

- 2 or more tasks **seem** to be performed simultaneously
- Implemented in most modern PCs
- How does it work?
- Recall: Scheduling based on Time-slicing and pre-emption



Parallel Execution

- 2 or more tasks **are** performed simultaneously
- Requires cooperation from hardware
 - Multiple CPUs
 - CPU with multicores
 - Multithreaded CPU



Threads

User Threads

- Thread management done by user-level threads library
- Three primary thread libraries:
 - POSIX **Pthreads**
 - Win32 threads
 - Java threads


Kernel Threads

- Supported by the Kernel
- Examples
 - Windows XP/2000
 - Solaris
 - Linux
 - Tru64 UNIX
 - Mac OS X

Question: What kind of relationship exists between User Threads and Kernel Threads?


Many-to-One Model

- **M:1 Model - Many** user threads mapped to a kernel thread
- Thread management done by thread library in user space
 - As many user threads can be created
- Only one thread can access the kernel
 - No support for parallel execution on multiprocessors
 - Process is blocked when one thread makes a system call




One-to-One Model

- **1:1-Model** - **One** user thread mapped to **one** kernel thread
- Pros: better support for concurrency
 - Process is NOT blocked when one thread makes a system call
 - Multiple threads run in parallel on multiprocessor
- Cons: requires a kernel thread for each new user thread
 - Degrade performance
 - # user threads can be limited



Many-to-Many Model

- **M:M – Model - Many** user threads mapped to **many** kernel thread
 - # kernel threads \leq # user threads
 - No limit on the # threads
- Support for concurrency
- *Two-level* model supports mapping one user thread to one kernel thread



Multicore Programming

- Multicore systems putting pressure on programmers, challenges include:
 - **Dividing activities**
 - Separate tasks
 - **Balance**
 - Tasks have equal work
 - **Data splitting**
 - Separate data
 - **Data dependency**
 - Ensure synchronization
 - **Testing and debugging**