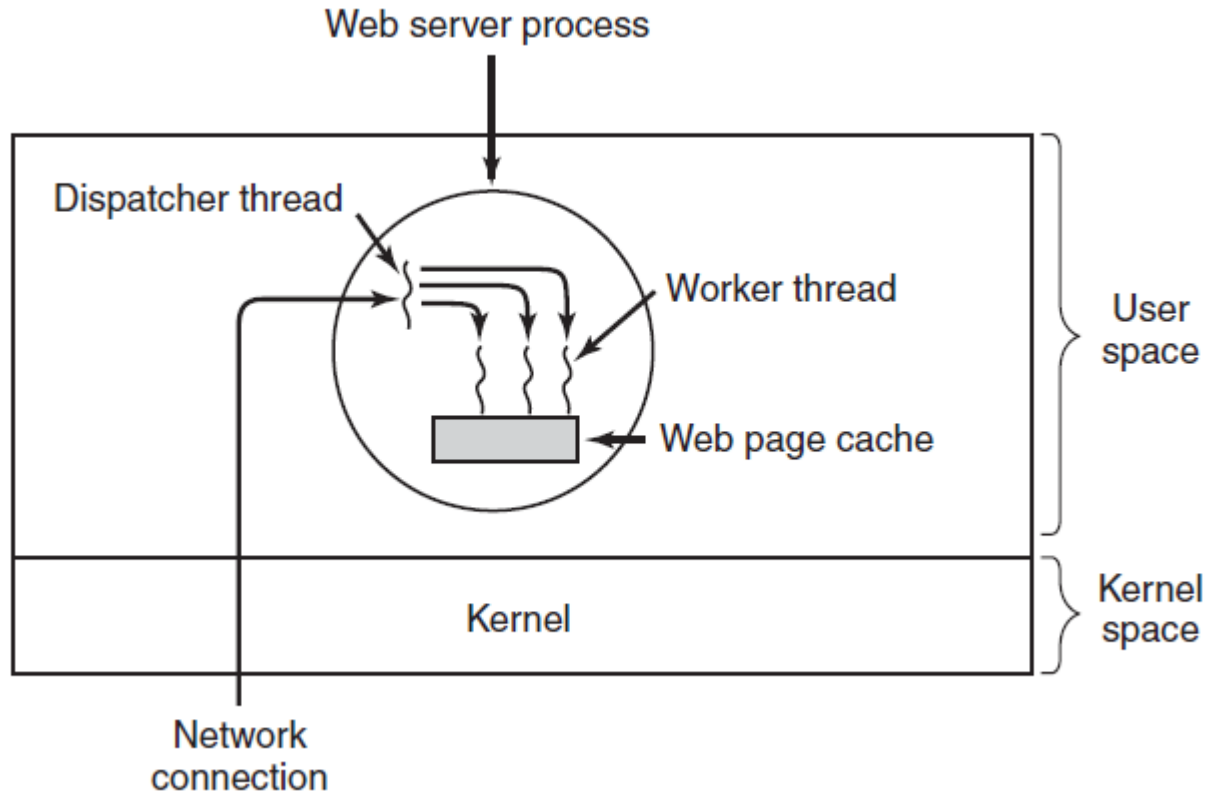


Agenda


- Assignment 1 due September 29!
- Today's lecture
 - Multiprogramming in Java
 - IPC
 - Issues with Concurrency
 - Synchronization
- Reading: Sections 3.4, 4.4-4.5, 6.1-6.2
- Next : More on Synchronization

Multiprogramming Examples

- Web Server





- Word processor



Thread Libraries

- Two primary ways of implementing threads
 - Library entirely in user space
 - Kernel-level library supported by the OS
- **Provide** programmer with API for creating and managing threads
 - POSIX Pthreads
 - Win32
 - Java



Pthreads

- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Solaris, Linux, Mac OS X)

Win32 Threads

- Kernel level library
- Shared data declared globally as with Pthreads

Java Threads

- Java threads are managed by the JVM
 - Implemented using host thread library
 - Win32 API on Windows
 - Pthreads on Unix and Linux
- User thread – kernel thread mapping
 - Depends on the host operating system
 - One-to-one model on Windows XP
- Java threads may be created by:
 - Implementing the Runnable interface

```
public interface Runnable
{
    public abstract void run();
}
```

Java Threads – Example Program

```
class Sum
{
    private int sum;

    public int get() {
        return sum;
    }

    public void set(int sum){
        this.sum = sum;
    }
}
```

```
class Summation implements Runnable
{
    private int upper;
    private Sum sumValue;

    public Summation(int upper, Sum sumValue)
    {
        this.upper = upper;
        this.sumValue = sumValue;
    }

    public void run() {
        int sum = 0;

        for (int i = 0; i <= upper; i++)
            sum += i;

        sumValue.set(sum);
    }
}
```

Java Threads - Example Program

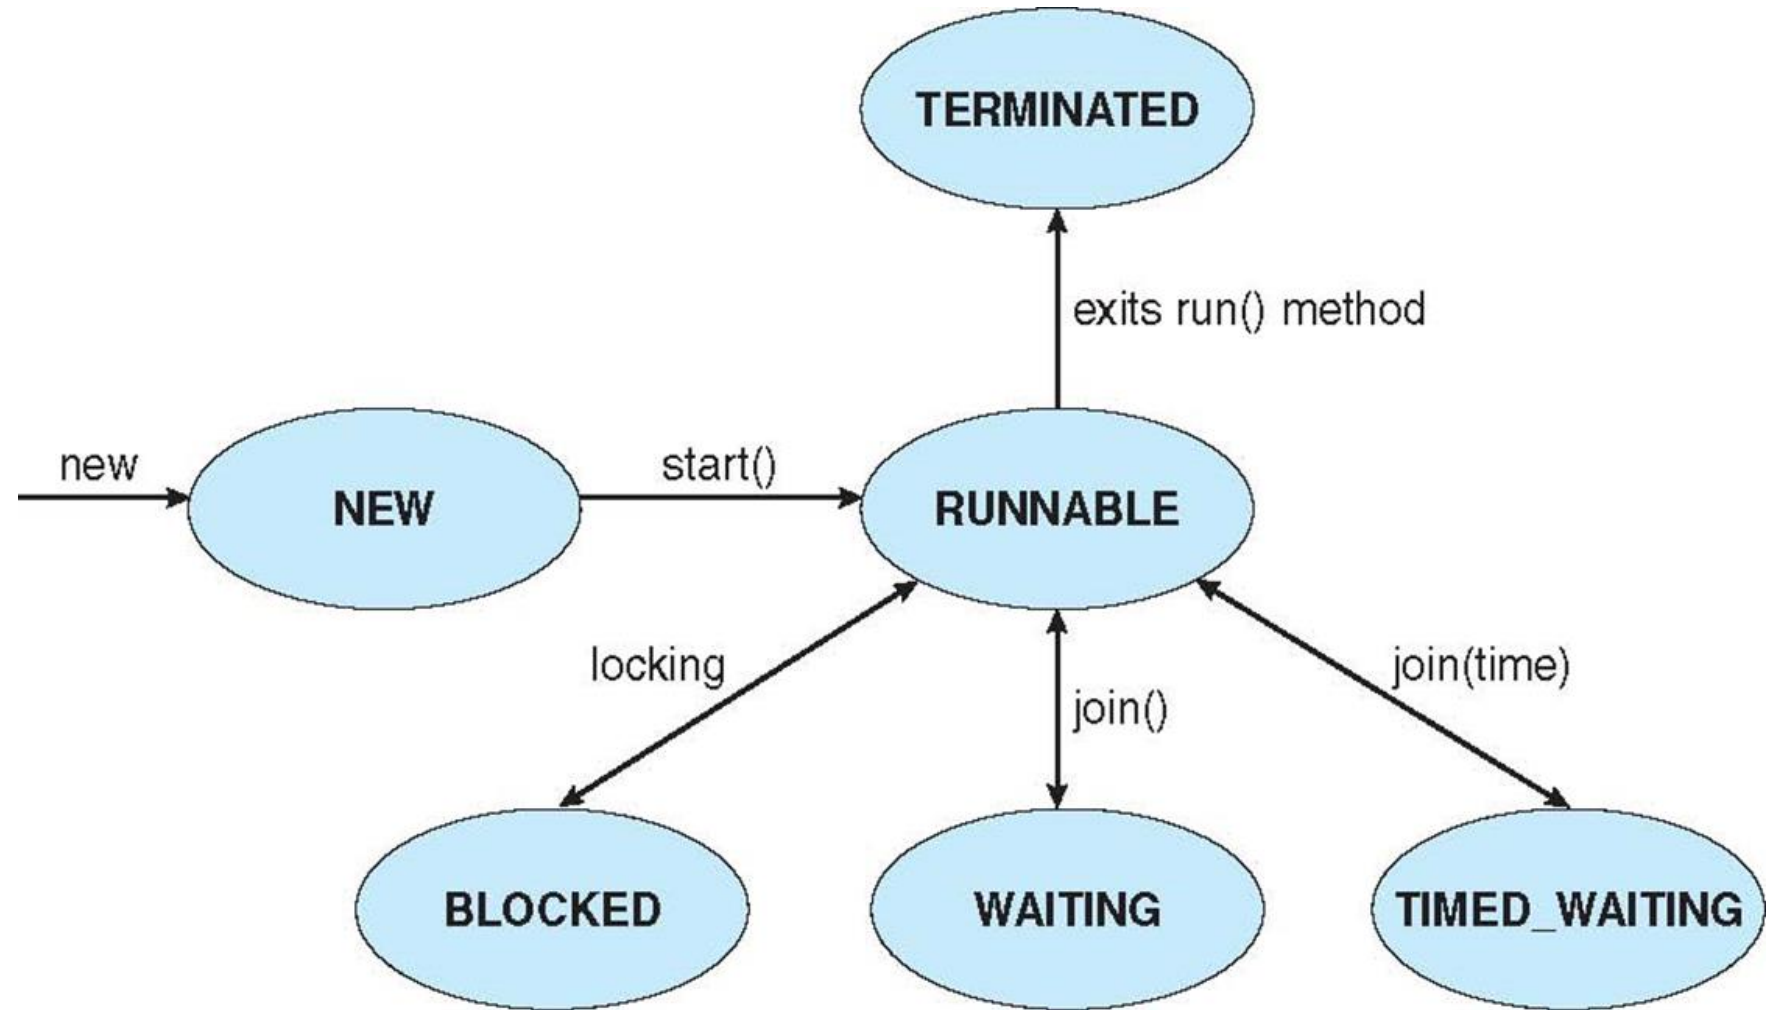
```
public static void main(String[] args) {  
    if (args.length != 1) {  
        System.err.println("Usage Driver <integer>");  
        System.exit(0);  
    }  
    if (Integer.parseInt(args[0]) < 0) {  
        System.err.println(args[0] + " must be >= 0");  
        System.exit(0);  
    }  
  
    // Create the shared object  
    Sum sumObject = new Sum();  
    int upper = Integer.parseInt(args[0]);  
  
    Thread worker = new Thread(new Summation(upper, sumObject));  
    worker.start();  
  
    try {  
        worker.join();  
    } catch (InterruptedException ie) { }  
    System.out.println("The sum of " + upper + " is " + sumObject.get());  
}
```

1. Allocate memory and initialize thread
2. Call the run () method

Wait for thread
termination

Shared data by passing reference to objects

Java Thread States



Threading Issues

- Semantics of **fork()** and **exec()** system calls
 - Does **fork()** duplicate only the calling thread or all threads?
- **Thread cancellation** of **target thread**
 - Asynchronous or deferred
- **Signal** handling
- **Thread pools**
- **Thread-specific data**
- **Scheduler activations**

Thread Cancellation

- Terminating a thread before it has finished
- Two general approaches:
 - **Asynchronous cancellation** terminates the target thread immediately
 - **Deferred cancellation** allows the target thread to periodically check if it should be cancelled
 - JAVA API provides the `Interrupt()` method which sets the interrupt status of the target thread
 - Threads periodically check their interrupt status
 - If interrupt status set, clean up and terminate

Signal Handling

- Signals are used in UNIX systems to notify a process that a particular event has occurred
- A **signal handler** is used to process signals.
 1. Signal is generated by particular event
 2. Signal is delivered to a process
 3. Signal is handled
- Options:
 - Deliver the signal to the thread to which the signal applies (e.g., synchronous signals)
 - Deliver the signal to every thread in the process (e.g., terminating signal <control><c>)
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process

Thread Pools

- Create a number of threads in a pool where they await work
- Advantages:
 - Usually slightly faster to service a request with an existing thread than create a new thread
- Allows the number of threads in the application(s) to be bound to the size of the pool
- `Java.util.concurrent` package includes an API for thread pools

Thread Specific Data

- Allows each thread to have its own copy of data
- Useful when you do not have control over the thread creation process (i.e., when using a thread pool)

Scheduler Activations

- Both M:M and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- Scheduler activations provide **upcalls** - a communication mechanism from the kernel to the thread library
- This communication allows an application to maintain the correct number of kernel threads


Concurrent Processes - Revisited

- Most OS provide process and thread abstractions to allow creation of multithread and multi-process application
- Hence, we talk about **concurrent processes**

Concurrent Processes - Revisited


Concurrent Execution

- 2 or more tasks **seem** to be performed simultaneously
- Implemented in most modern PCs
- How does it work?
- Recall: Scheduling based on Time-slicing and pre-emption




Parallel Execution

- 2 or more tasks **are** performed simultaneously
- Requires cooperation from hardware
 - Multiple CPUs
 - CPU with multicores
 - Multithreaded CPU



Concurrent Execution

- Multitasking, multiprocessing, or multiprogramming
 - Processes **seem** to be running simultaneously
- Achieved using Time slicing and Preemption
 - A process, Tx, gets to run for 1 time quantum (e.g., 10-100ms)
 - When quantum expires, Tx is interrupted
 - Scheduler selects Ty to run next
 - Tx is placed in the ready queue
 - Ty runs from where it was last interrupted
 - Repeat



Cooperating Processes

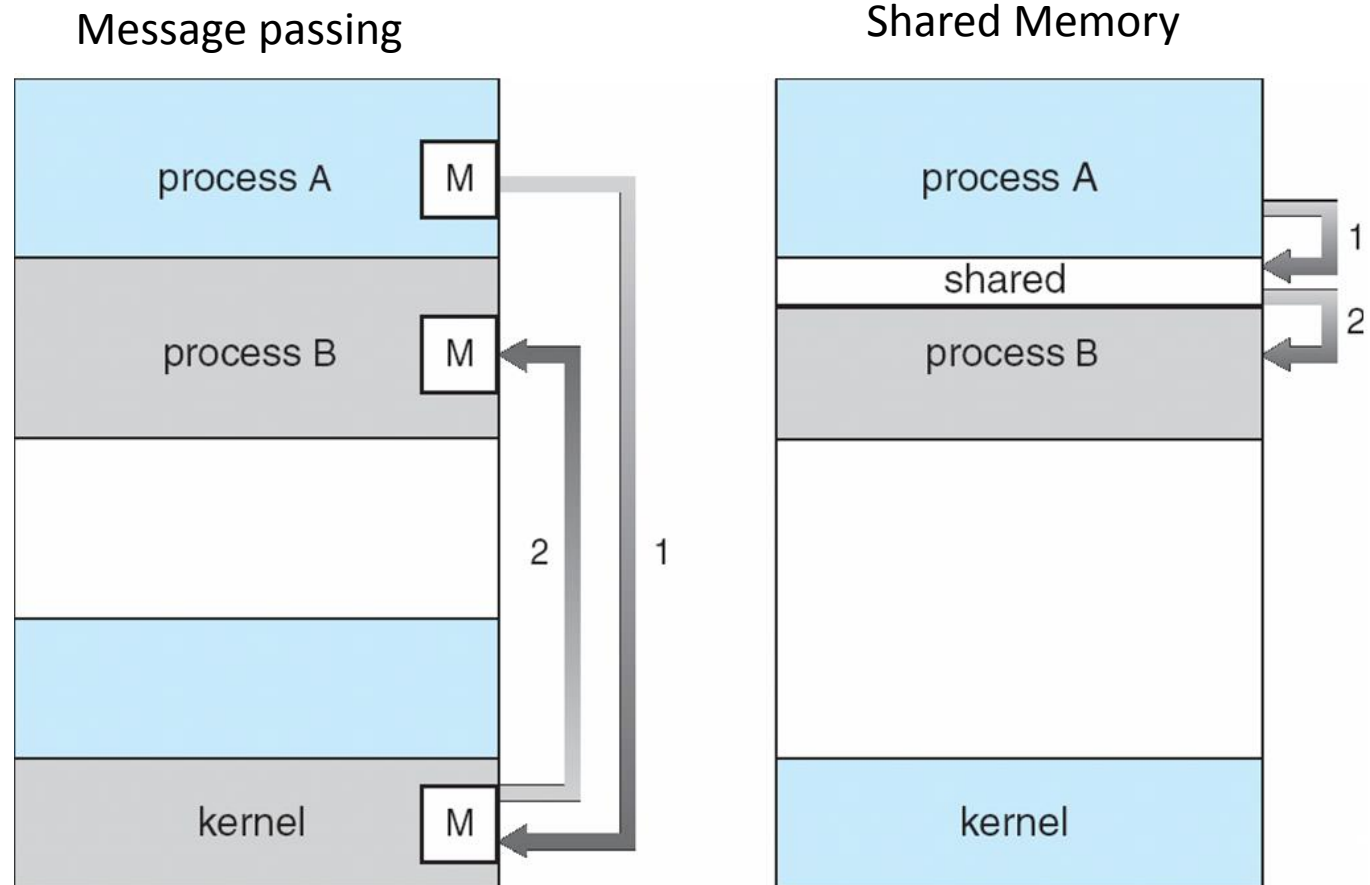
- Processes within a system may be **independent** or **cooperating**
- **Cooperating** process can affect or be affected by other processes, including sharing data
- Reasons for cooperating processes:
 - Information sharing: e.g., files
 - Computation speedup: e.g., distribute workload among CPUs
 - Modularity: e.g. client/server architecture
 - Convenience: e.g., facilitate maintenance, security

Cooperating Processes

- Cooperating processes need
 - **Interprocess communication (IPC)**: exchanging data from one process to another
 - Synchronization: coordinating process activities
 - All threads of a process share the same address space and other resources
 - Any alteration of a resource by one thread affects the other threads in the same process
- The OS typically provides support for both IPC and synchronization

Interprocess Communication

- Cooperating processes need **interprocess communication (IPC)**
- Two models of IPC
 - Shared memory
 - E.g., global variables
 - E.g., threads
 - Message passing
 - E.g., e-mail, TCP/IP



Difficulties of Concurrency

- **By Definition, a cooperating** process can affect or be affected by the execution of another process
 - Recall: OS tries to prevent one process from accessing another process's memory
 - But IPC uses shared memory!!
- Concurrent access to shared data may result in data inconsistency
 - Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes
- Difficult for the OS to manage the allocation of resources optimally
- Difficult to locate programming errors as results are not deterministic and reproducible

Producer-Consumer Problem

- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
 - Word processor produces files to be consumed by the printer
 - Client/server paradigm
 - Server: web server produces files
 - Client: web browser consumes files


Producer-Consumer Problem

- The Producer-Consumer Problem consists of:
 - 2 or more processes or threads
 - 1 or more processes produce data
 - 1 or more processes consume the data
 - A shared buffer (or queue)
 - Producer(s) must safely pass the data to consumer(s)
- Two approaches
 - *unbounded-buffer*: places no practical limit on the size of the buffer
 - Producer can produce at will
 - Consumer can consume only if the buffer is not empty
 - *bounded-buffer*: assumes that there is a fixed buffer size
 - Producer can produce only if the buffer is not full
 - Consumer can consume only if the buffer is not empty

Producer-Consumer Problem

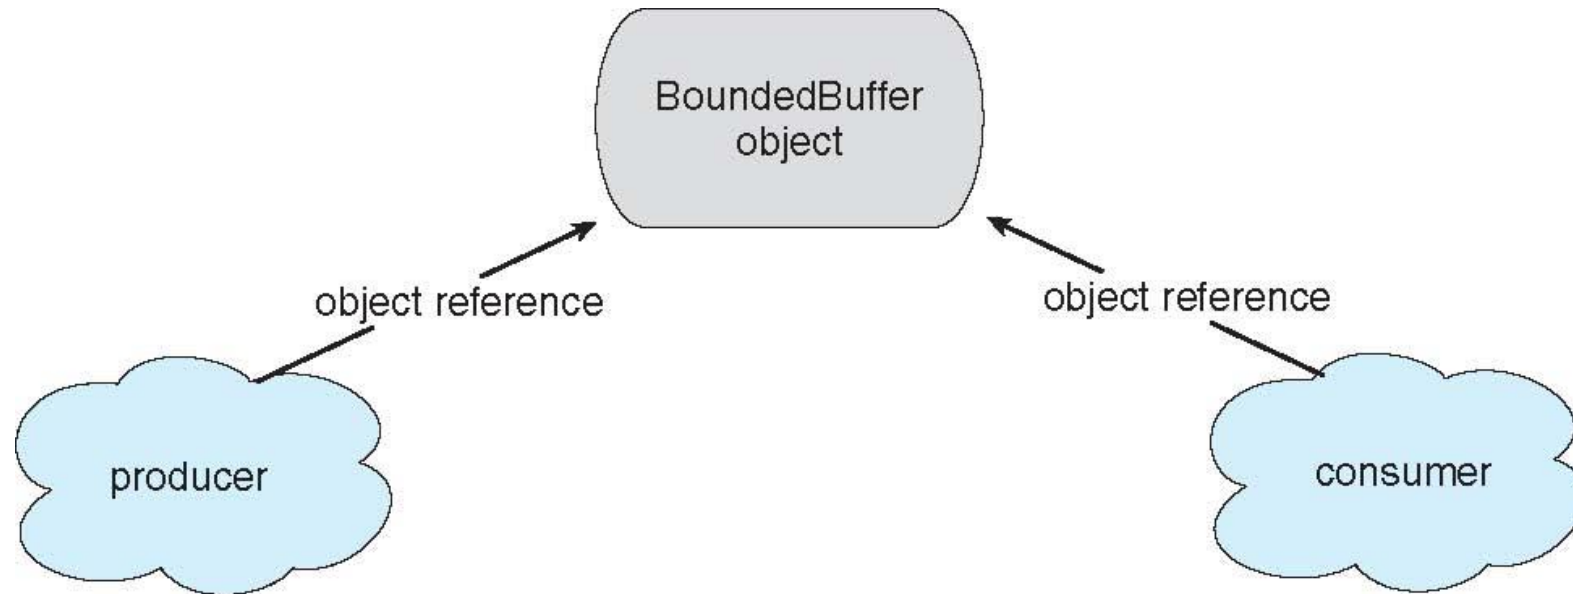
- Synchronization Challenges:

- Empty buffer
- Full buffer
- Shared variables



Simplified Producer-Consumer

Simulating Shared memory in Java



Java Threads – Producer - Consumer

```
public interface Buffer <E>
{
    // producers call this method
    public void insert(E item);

    // consumers call this method
    public E remove();
}
```

```
public class BufferImpl<E> implements Buffer<E>{
    private static final int BUFFER_SIZE = 5;
    private E[] elements;
    private int in, out, count;

    public BufferImpl() {
        count = 0;
        in = 0;
        out = 0;
        elements = (E[]) new Object[BUFFER_SIZE];
    }


    public void insert(E item) {} // producers call this method
    public E remove() {} // consumers call this method
}
```

Java Threads – Producer - Consumer

// producers call this method

```
public void insert(E item) {
```

Wait until buffer
is not full



```
while (count == BUFFER_SIZE)
```

```
; // do nothing -- no free space
```

```
// add an element to the buffer
```

```
elements[in] = item;
```

```
in = (in + 1) % BUFFER_SIZE;
```

```
++count;
```


```
}
```

// consumers call this method

```
public E remove() {
```

```
E item;
```

Wait until buffer
is not empty



```
while (count == 0)
```

```
; // do nothing - nothing to consume
```

```
// remove an item from the buffer
```

```
item = elements[out];
```

```
out = (out + 1) % BUFFER_SIZE;
```

```
--count;
```

```
return item;
```


```
}
```

Producer - Consumer

```
public static void main(String[] args) {  
    // create the shared object  
    Buffer<String> boundedBuffer = new BufferImpl<String>();  
  
    // create the producer and consumer threads  
    Thread producer = new Thread(new Producer(boundedBuffer));  
    Thread consumer = new Thread(new Consumer(boundedBuffer));  
  
    // start the threads  
    producer.start();  
    consumer.start();  
}
```

Simplified Producer-Consumer

Synchronization Challenge: what is the value of **count** assuming producer and consumer threads run concurrently?



Synchronization – Shared Variable

- Machine instructions for ++count

```
move count → regA
inc regA
move regA → count
```
- Machine instructions for --count

```
move count → regB
dec regB
move regB → count
```
- The Producer and Consumer threads execute these instructions concurrently
- Problem: These instructions can be interleaved!

Interleaving Scenario 1

Instruction	count	regA	regB	
move count → regA	5	5	?	
inc regA	5	6	?	
move regA → count	6	6	?	Context switch
move count → regB	6	6	6	
dec regB	6	6	5	
move regB → count	5	6	5	

Correct: **count** has the same value before and after an increment and decrement!

Interleaving Scenario 2

Instruction	count	regA	regB	
move count → regA	5	5	?	Context switch
move count → regB	5	5	5	
dec regB	5	5	4	
move regB → count	4	5	4	Context switch
Inc regA	4	6	4	
move regA → count	6	6	4	

Incorrect: **count** has the value as if decrement never happened!

Interleaving Scenario 3


Instruction				count	regA	regB	
move count → regB				5	?	5	
dec regB				5	?	4	Context switch
move count → regA				5	5	4	
inc regA				5	6	4	
move regA → count				6	6	4	Context switch
move regB → count				4	6	4	

Incorrect: **count** has the same value as if increment never happened!

Thread Scheduling

- What was happening?
 - The same code is executed in each scenario
 - The threads were scheduled differently
 - Both threads write to `count`
 - Different behavior depending on when the writes occur
- **BUT**, thread scheduling can be non-deterministic!
- Scheduling in scenarios 2 and 3 resulted in **destructive interference**

Race Condition



- Two processes want to access shared memory at the same time
- A race condition is when *destructive interference* can occur.
- With increasing parallelism due to increasing number of cores, race condition are becoming more common

Critical Region (Section)

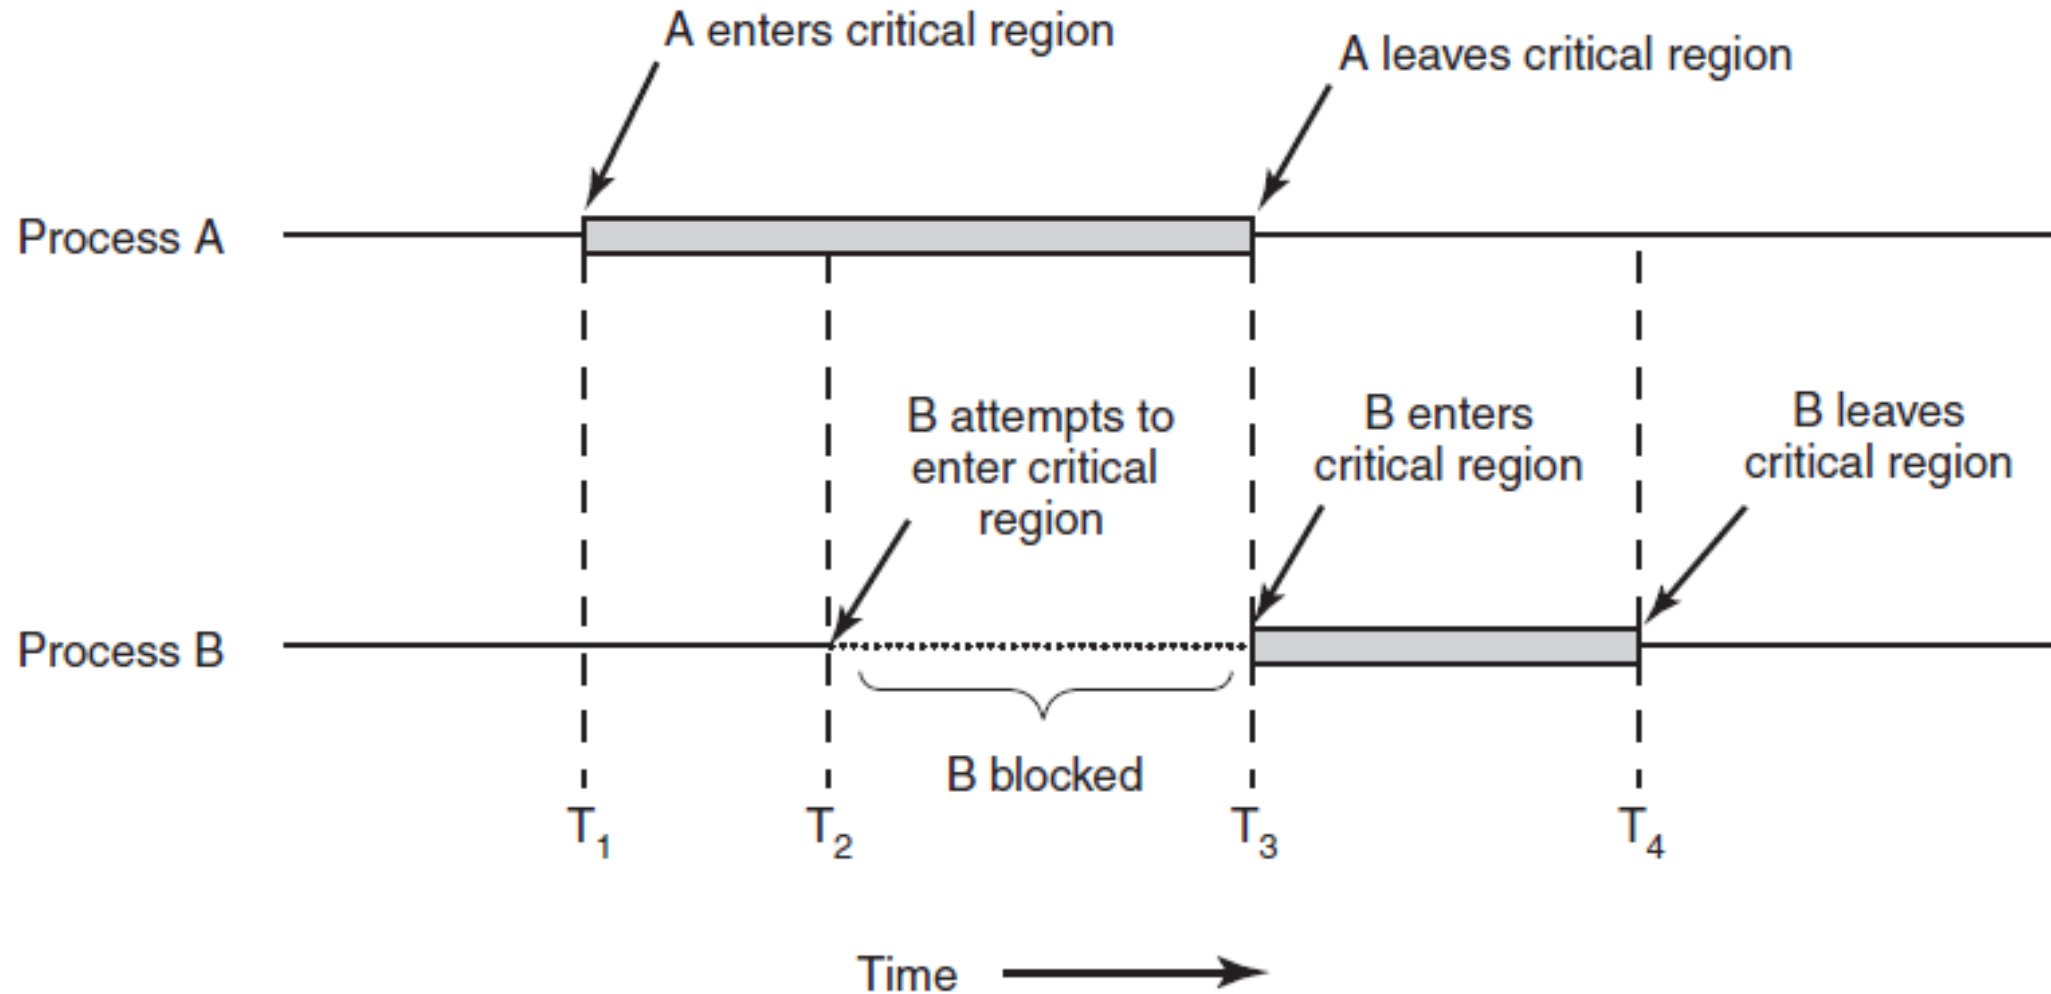
- Part of the program where shared memory is accessed
- Prone to race conditions
- **Problem:** Multiple threads executing concurrently in a critical section can cause errors
- How do we prevent this from happening?

Critical Regions

Requirements to avoid race conditions:

1. **Mutual exclusion:** No two processes may be simultaneously inside their critical regions.
2. **Scheduler independent:** No assumptions may be made about speeds or the number of CPUs.
3. **Allows progress:** No process running outside its critical region may block other processes.
4. **Starvation free:** No process should have to wait forever to enter its critical region.

Mutual Exclusion Using Critical Region



Next...

- Solutions to synchronization issues
 - Software solutions
 - Hardware solutions