

Agenda

- Assignment 3 due Friday November 10
- Today's lecture
 - IPC
 - Message Passing
 - Pipes
 - Signals
 - Memory management
 - Textbook Reading: 8.1-8.3
 - Next Lecture: Partitioning, Paging, Segmentation

Where we are...

- The OS provides three key abstractions
 - Process → CPU
 - Memory (Address) Space → RAM
 - Files → Secondary storage, Network, and Peripheral devices

Message Passing

- When processes interact with one another two fundamental requirements must be satisfied:
 - Synchronization
 - Communication
- Message passing is one approach to providing both of these functions
 - Works with distributed systems *and* shared memory multiprocessor and uniprocessor systems
- Features:
 - Post-office or telephone model
 - Processes cannot read each others' memory
 - To communicate, processes ask the system to send a message on their behalf
 - Real life example: the Internet

Message Passing Operation

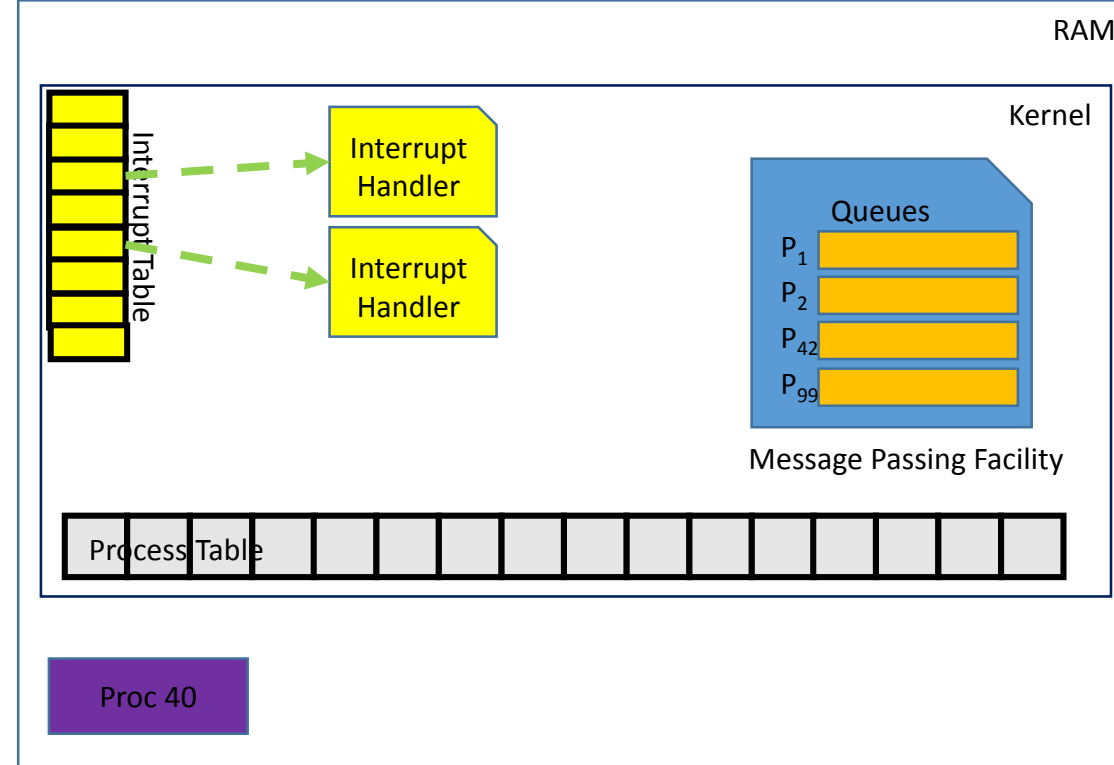
- The OS provides a “Post-office” facility
- The actual function is normally provided in the form of a pair of primitives:

```
send (destination, message)
receive (source, message)
```

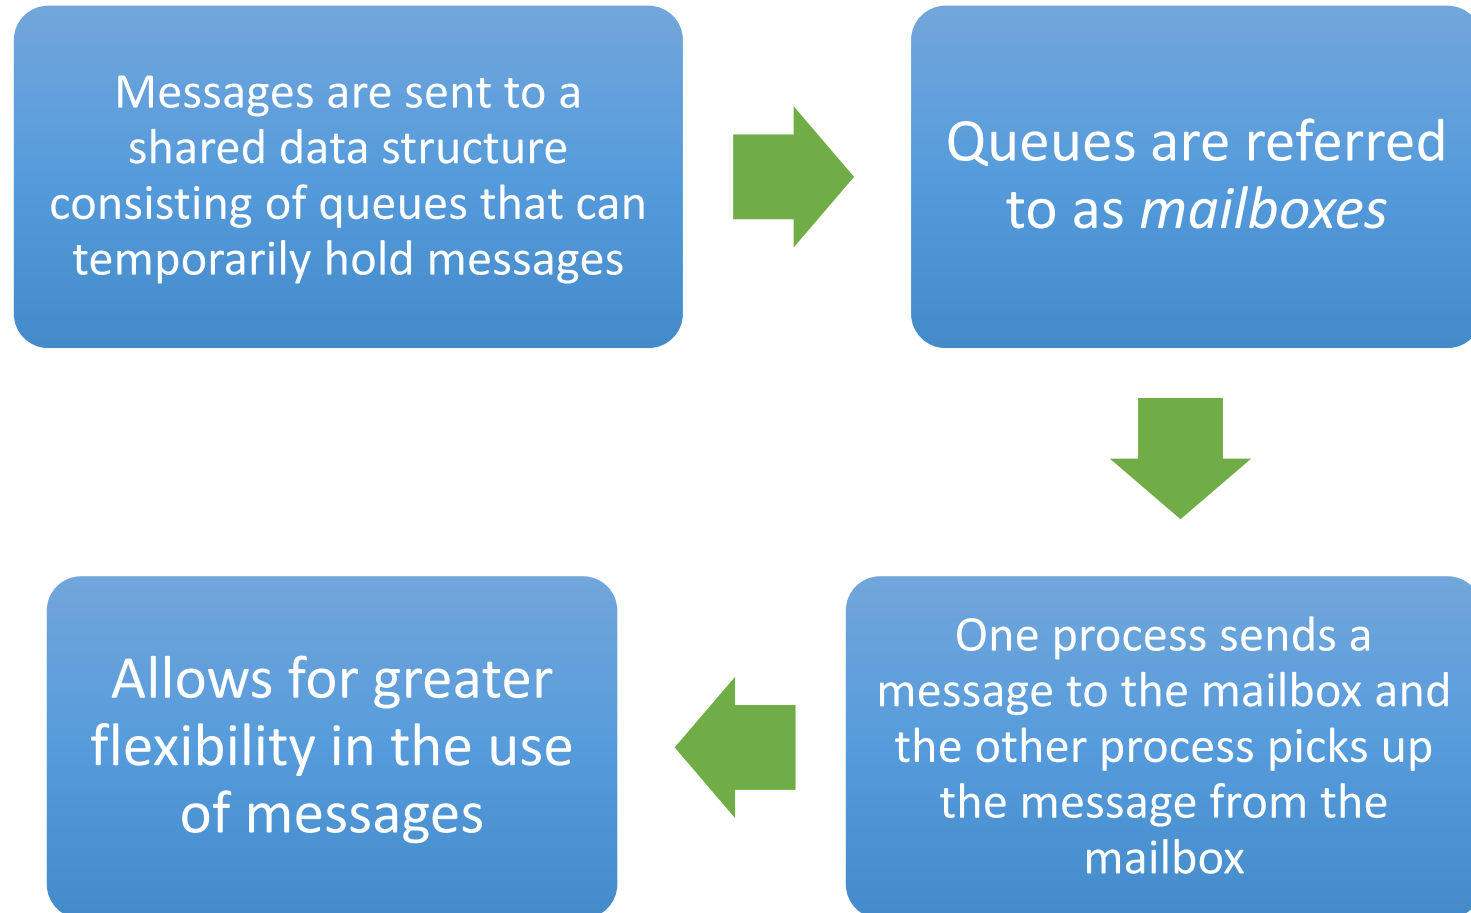
- A process sends information in the form of a *message* to another process designated by a *destination*
- A process receives information by executing the `receive` primitive, indicating the *source* and the *message*

The Message Passing Facility

- Queue P_i stores messages (or senders) to the process P_i
- The head of queue P_i is the next message to be received by P_i
- Sending processes:
 - must know the ID or name of the process they are sending to
 - ask the kernel to enqueue the message on the corresponding queue
- To receive a message, it is not (always) necessary for the process to explicitly specify the receiver
 - The process asks the kernel to remove a message from their queue and return it
- Analogy: A queue is the same as a mailbox




Message Passing – Indirect Addressing




Sender Receiver Relationship




(a) One to one



(b) Many to one




(c) One to many



(d) Many to many

Message Format

- The message is divided into two parts:
 - **header**, which contains information about the message
 - an identification of the source and intended destination of the message,
 - a length field, and a type field to discriminate among various types of messages
 - control information
 - **body**, which contains the actual contents of the message



Process Naming

- How processes are named and how process names are discovered is a whole topic beyond scope of this course

Example: On the Internet

- Processes are named by IP/port #
- Use DNS to look up IP addresses and port #s of specific processes: e.g., 25 for e-mail (SMTP)
- Other Examples:
 - On a single system each process has unique process ID (pid) assigned by the OS
 - Each thread has a unique thread ID (tid), assigned by the OS or run-time system such as Java
- Another idea is to use a registry (e.g., DNS)

Synchronization


Communication of a message between two processes implies synchronization between the two

The receiver cannot receive a message until it has been sent by another process

When a receive primitive is executed in a process there are two possibilities:

If there is no waiting message the process is blocked until a message arrives or the process continues to execute, abandoning the attempt to receive

If a message has previously been sent the message is received and execution continues



Should Processes Block?

- Questions:
 - When a process sends a message, should it block until the message is received?
 - When a process receives a message, should it block until the message arrives?
- Answers depend on the system:
 - Typically, receivers will block until the message arrives
 - Whether senders should block depends on what type of behaviour is desirable:
 - Synchronous sends
 - Asynchronous sends

Blocking Send, Blocking Receive

- Both sender and receiver are blocked until the message is delivered
- Sometimes referred to as a *rendezvous*
- Allows for tight synchronization between processes
- At end of operation both sender and receiver have synchronized
i.e., Both know where they are
- Analogy: Telephone

Nonblocking Send

Nonblocking send, blocking receive

- Sender continues on but receiver is blocked until the requested message arrives
- Most useful combination
- Sends one or more messages to a variety of destinations as quickly as possible
- Example -- a service process that exists to provide a service or resource to other processes

Nonblocking send, nonblocking receive

- Neither party is required to wait
- Sender does not know when receiver gets the message
- Analogy: sending a letter

Producer Consumer Solution with Message Passing

```
void Producer() {  
    while( 1 ) {  
        p = produce();  
        Send( consumer, p );  
    }  
}
```

```
void Consumer( ) {  
    while( 1 ) {  
        Recv( producer, p )  
        consume( p );  
    }  
}
```

Pipes

- Special form of message passing
- Processes establish a direct connection (pipe) between themselves
- Can pass messages or data streams to each other.
- Pipes can either be
 - Local: both processes on the same machine
 - Interhost: (established via a network connection)
- Processes need to be able to identify each other



Signals

- Processes can asynchronously notify other processes
- All processes running on the same machine
- All processes controlled by the same OS
- One process asks the OS to notify another process on its behalf

Memory Management

Objectives

- Discuss requirements for memory management
- Various memory-management techniques
 - Partitioning
 - Paging
 - Segmentation

Motivation


- Main memory and registers are only storage CPU can access directly
 - Register access in one CPU clock (or less)
 - Main memory can take many cycles
 - **Cache** sits between main memory and CPU registers
- Memory is allocated when
 - A process is created
 - A process needs more heap/stack
 - The OS needs to use cache
 - The OS needs to allocate data structures

Motivation

- Memory is a finite resource and needs to be managed
 - Fair allocation among processes (sharing)
 - Efficient allocation --> OS has enough memory
 - Protection of memory required to ensure correct operation
- Issues:
 - Program may not be at address 0x00000000h
 - Linking and loading issues
 - Virtual memory may be used
 - Mapping and paging issues
 - Caching issues
 - Components may be duplicated
- How does the OS handle these situations?

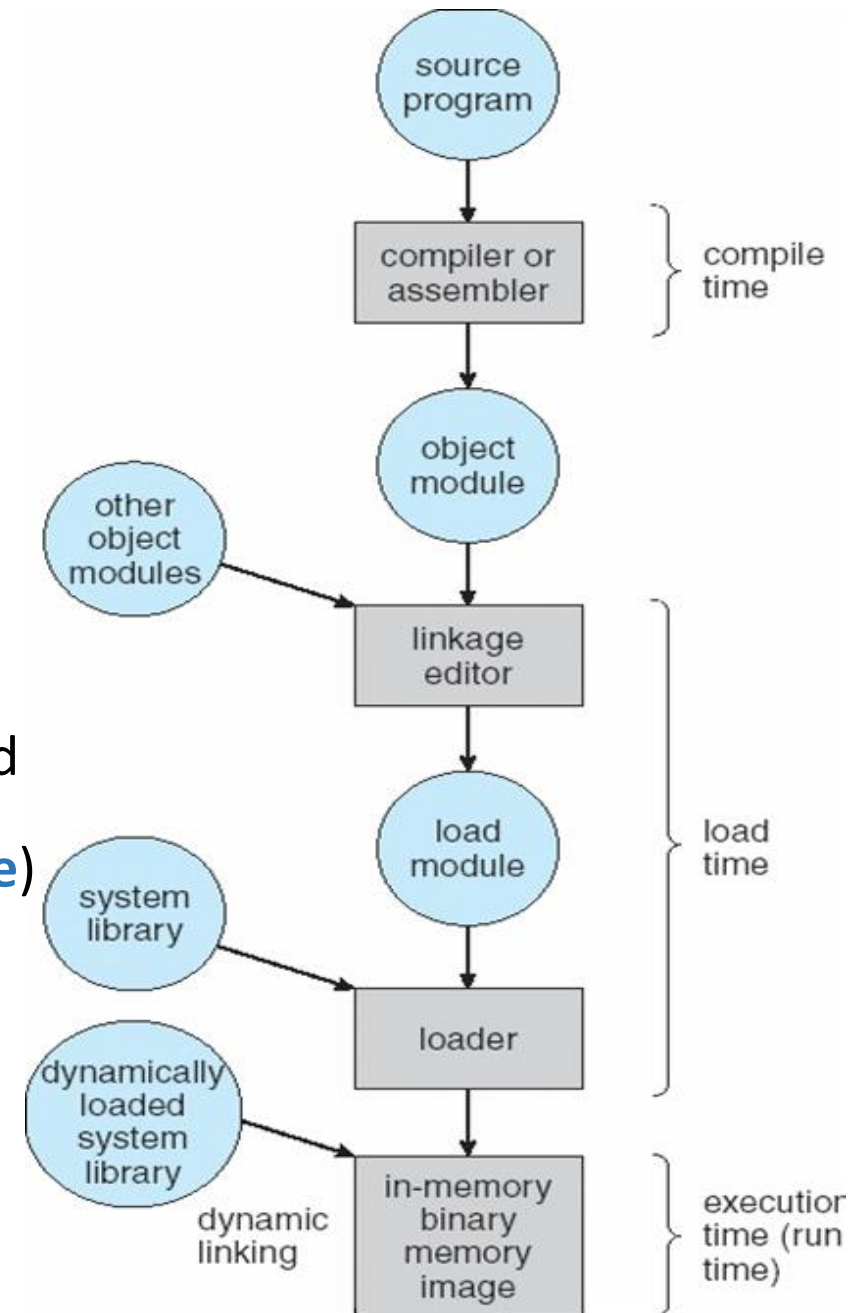
What is Memory?

- Computer memory consists of a linear array of addressable storage cells that are similar to registers
- Program must be brought (from disk) into memory and placed within a process for it to be run
 - Each process has a base/limit address to specify its accessible addresses
 - "segmentation fault" error generated when trying to access memory outside of the allowable address space



Address Binding

- Linking a symbolic address in a program to an actual memory location in main memory
- **Compile time**
 - The compiler puts physical addresses in the program (**absolute code**)
- **Loading time**
 - Compiler uses offsets from some base which is determined at load time
 - The addresses are inserted by the loader (**relocatable code**)
- **Execution time**
 - Need hardware support (base and limit registers) to map addresses
- Most modern Oss use virtual addressing instead



Memory Management Goals

- Make sure each process has sufficient memory
- Keep as many processes in memory as possible
- Allocate memory efficiently
 - Allocate as much as is needed for a process – not more
 - Leave some free space for new starting processes
 - Maximize memory utilization
- Memory is a finite resources
- These goals cannot be simultaneously met

OS Strategies

- Swapping
- Segmentation
- Paging
- Virtual memory

We will discuss those strategies in the next few lectures.