

# Agenda

- Assignment 2 is out, due October 20!
- Exam Schedule is out!
- Today's lecture
  - Critical section problem
  - Software and hardware solution
- Reading: Sections 3.4, 4.4-4.5, 6.1-6.2
- Next : More on Synchronization

# Mutual Exclusion

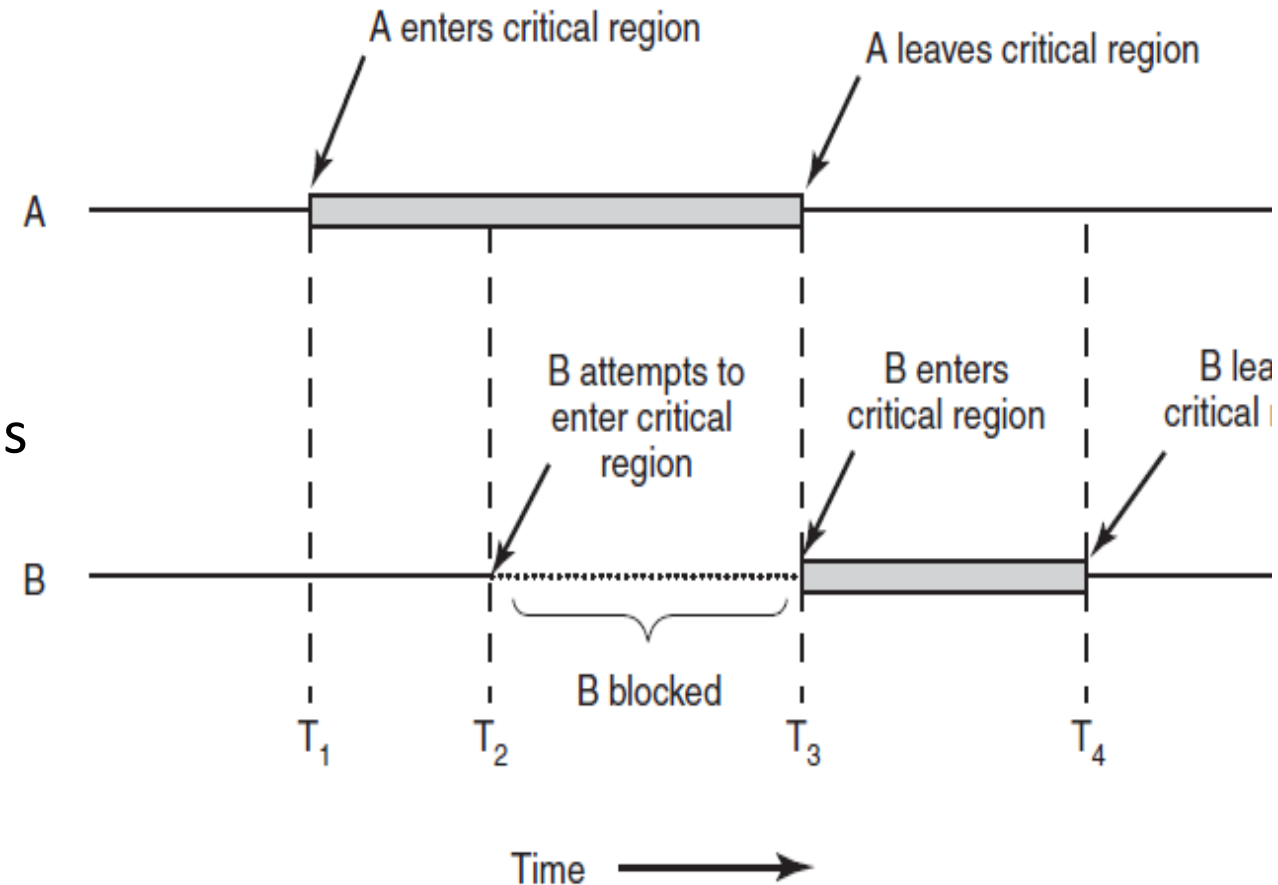
- Consider a block of shared memory
- Process P wants to write data to the shared memory while Process R wants to read the data
- **Sequential execution:** a synchronous system call used to cause the R to wait until the P is complete
  - The approach works because the synchronous system calls provide mutual exclusion
- **Concurrent execution:** maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes

# Mutual Exclusion

- **Destructive Interference:** when cooperation among concurrent processes results in incorrect behavior
- **Critical section** (region): section of code within a process that requires access to shared resources and must not be executed while another process is in a corresponding section of code
- **Race Condition:** A case where multiple threads read and write a shared data item and the final result depends on the relative timing of their execution
- When a resource (i.e., memory) can be accessed by only one process at a time, we say that the resource is **mutually exclusive**
- To prevent race conditions, we use mutual exclusion and critical sections

# Critical Region Requirements

1. **Mutual exclusion:** No two processes may be simultaneously inside their critical regions.
2. **Scheduler independent:** No assumptions may be made about speeds or the number of CPUs.
3. **Allows progress:** No process running outside its critical region may block other processes.
4. **Starvation free:** No process should have to wait forever to enter its critical region



# Typical Process Structure

```
while (true) {  
    entry section  
    Critical section  
    exit section  
    remainder section  
}
```

# Locks

- Locks can be used as mutual exclusion mechanisms
  - Prevent threads from entering a critical section if another thread is present
  - Allow threads to wait, and eventually enter their critical section
- Locks support two operations
  - Acquire: locks the critical section so that it is safe to enter
  - Release: unlock the critical section on exit
- Both hardware and software solutions are possible

```
while (true) {  
    acquire lock  
        critical section  
    Release lock  
        remainder section  
}
```

# Software Solution 1: Lock Variable

```
// thread
while (true) {
    entry section
    while (busy);
    busy = true
    critical section
    busy = false
}
```

Busy waiting

busy is a shared variable

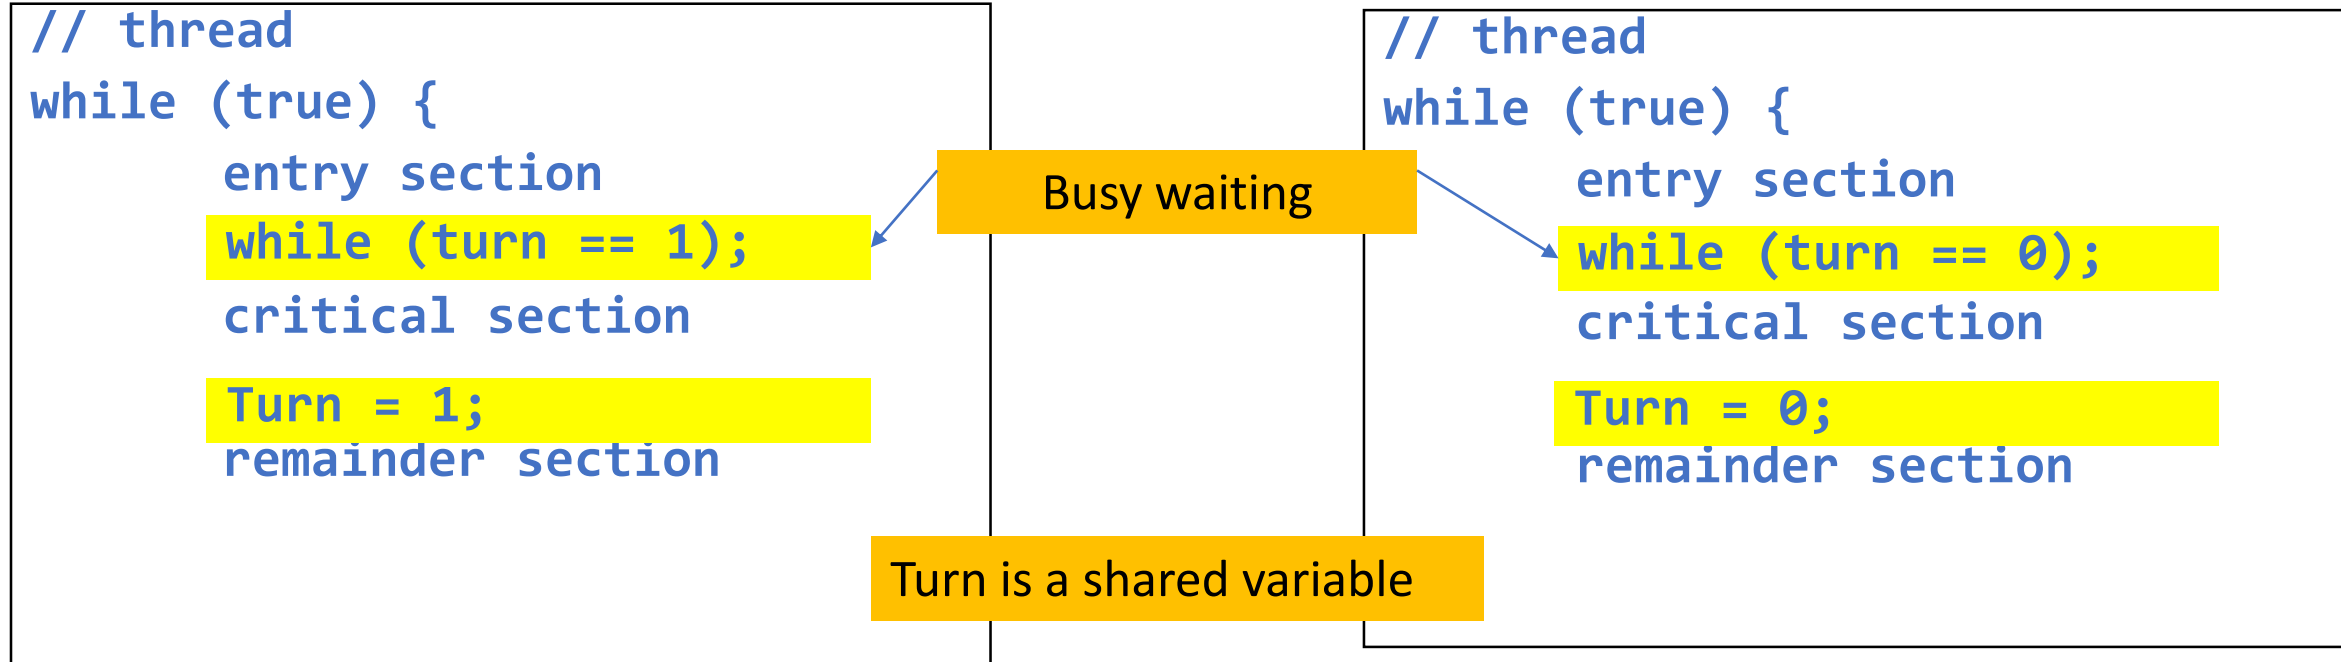
```
// thread
while (true) {
    entry section
    while (busy);
    busy = true
    critical section
    busy = false
}
```

1. Mutual exclusion: 😞
2. Scheduler independent: 😊
3. Allows progress: 😊
4. Starvation free: 😞

## Problems

- Requires busy waiting
- Prone to starvation

# Software Solution 2: Strict Alternation



1. Mutual exclusion: 😊
2. Scheduler independent: 😊
3. Allows progress: 😞
4. Starvation free: 😊

## Problems

- Works for 2 processes!
- Requires busy waiting
- One thread may block the other!



# Software Solution 3: Peterson's Solution

- Based on Dekker's solution
- Processes share two variables:
  - int **turn**;
  - boolean **flag[2]**
- **Strict alternation**: the variable **turn** indicates whose turn it is to enter the critical section.
- **Lock**: the **flag** array is used to indicate if a process is ready to enter the critical section. **flag[i]** = true implies that process **P<sub>i</sub>** is ready!

# Software Solution 3: Peterson's Solution

//Producer thread

```
while (true) {  
    entry section  
    flag[i] = true;  
    turn = j;  
    While (flag[j] && turn == j);  
    critical section  
    flag[i] = false;  
}
```

Entering CS

Let other thread  
go first.

flag and turn are shared

//Consumer thread

```
while (true) {  
    entry section  
    flag[j] = true;  
    turn = i;  
    While (flag[i] && turn == i);  
    critical section  
    flag[j] = false;  
}
```

Wait while busy

Leaving CS

1. Mutual exclusion: 😊
2. Scheduler independent: 😊
3. Allows progress: 😊
4. Starvation free: 😊

## Problems

- Works for 2 processes! 😞
- Requires busy waiting 😞
- Assumes writes and reads are atomic! 😞

# Hardware Solution: Interrupt Disabling

- Many systems provide hardware support for critical section code
- **Uniprocessors – could disable interrupts**
- In a uniprocessor system, concurrent processes cannot have overlapped execution; they can only be interleaved
- A process will continue to run until it invokes an OS service or until it is interrupted
- Therefore, to guarantee mutual exclusion, it is sufficient to prevent a process from being interrupted
- This capability can be provided in the form of primitives defined by the OS kernel for disabling and enabling interrupts

# Hardware Solution: Interrupt Disabling

## Disadvantages

- The efficiency of execution could be noticeably degraded because the processor is limited in its ability to interleave processes
- This approach will not work in a multiprocessor architecture
- Modern machines provide special *atomic* hardware instructions

# Atomic Instruction

- Atomic Instruction: a function or action implemented as a sequence of one or more instructions that appears to be indivisible.
  - No other process can see an intermediate state or interrupt the instruction
  - The sequence of instruction is guaranteed to execute as a group, or not execute at all
  - Atomicity guarantees isolation from concurrent processes

# Hardware Solution: Special Machine Instruction

## Test&Set Instruction (TSL instruction)

- Takes one operand, a Lock variable, to coordinate access to shared memory
- A **test** is made between memory value and a test value
- If the values are the same, a **set** value occurs

enter\_region:

TSL REGISTER,LOCK

| copy lock to register and set lock to 1

CMP REGISTER,#0

| was lock zero?

JNE enter\_region

| if it was nonzero, lock was set, so loop

RET

| return to caller; critical region entered

leave\_region:

MOVE LOCK,#0

| store a 0 in lock

RET

| return to caller

# Hardware Support: Using Test and Set

```
// thread
while (true) {
    entry section
    while (TestSet(busy));
    critical section
    busy = false
}
```

1. Mutual exclusion: 😊
2. Scheduler independent: 😊
3. Allows progress: 😊
4. Starvation free: 😞

## Problems

- Requires busy waiting 😞
- Starvation is possible

# Hardware Solution: Special Machine Instruction

## Compare&Swap Instruction (XCHG instruction)

- Also called a “compare and exchange instruction”
- A **compare** is made between a memory value and a test value
- If the values are the same a **swap** occurs
- Carried out atomically (not subject to interruption)

enter\_region:

```
MOVE REGISTER,#1
XCHG REGISTER,LOCK
CMP REGISTER,#0
JNE enter_region
RET
```

```
| put a 1 in the register
| swap the contents of the register and lock variable
| was lock zero?
| if it was non zero, lock was set, so loop
| return to caller; critical region entered
```

leave\_region:

```
MOVE LOCK,#0
RET
```

```
| store a 0 in lock
| return to caller
```



# Hardware Solution: Special Machine Instruction

## Advantages

- It is applicable to any number of processes on either a single processor or multiple processors sharing main memory
- It is simple and therefore easy to verify
- It can be used to support multiple critical sections; each critical section can be defined by its own variable

# Hardware Solution: Special Machine Instruction

## Disadvantages

- **Busy waiting is employed**
  - Thus, while a process is waiting for access to a critical section, it continues to consume processor time
- **Starvation is possible**
  - When a process leaves a critical section and more than one process is waiting, the selection of a waiting process is arbitrary. Thus, some process could indefinitely be denied access
- **Deadlock is possible**
- **Difficult to implement**

# Sleep and Wakeup

- Previous solutions require busy waiting!
  - Wasted CPU time
  - Unexpected effects (e.g., *priority inversion problem*)
- ***Sleep*** causes the caller to block, until another process wakes it up
- ***Wakeup*** causes a thread to resume work after it was put to sleep
- How is this helpful?

# Sleep and Wakeup

- To enter critical section (lock)
  - If thread can enter critical section
    - Enter critical section (acquire lock)
  - else
    - Add thread to a queue
    - Put thread to sleep until critical section is available
- When leaving critical section (unlock)
  - If another thread is waiting to enter critical section
    - Remove thread from queue
    - Wake thread
  - Leave critical section
- When a thread is woken (resumed)
  - Thread enters critical section (locks it)

```
while (true) {  
  
    acquire lock  
  
        critical section  
  
    Release lock  
  
        remainder section  
  
}
```

# Implementation Issues

- How do we lock the queue in order to add/remove threads?
  - Back to the original problem
- What happens if we resume a thread before it suspends itself?
  - In some systems it does not matter (In Java it does)
  - Thread can make sure that the resume “worked” (tricky)

# Alternative

Let the System be responsible for entering/leaving critical sections

- Advantages:
  - Since System is in charge, no race conditions in entering/leaving critical sections
  - Programmers do not have to write/debug critical section entry/leave routines
  - Code becomes simpler
- Disadvantages:
  - Code becomes system
  - These routines are more expensive, since they invoke the system
  - May do more than necessary
- In general a “System Solution” is a good thing
- Question: What kind of interface should the System provide?

# Abstractions for Mutual Exclusion

- We have several options:
  - Locks (mutex locks)
    - Usually not a good idea
    - Many implementations do not allow a thread to acquire its own lock
    - A spin lock is a lock mechanism that requires the thread to spin in a loop testing a condition
  - Semaphores
  - Monitors

# Semaphore

- Synchronization tool that does not require busy waiting
- Abstraction: Semaphore  $S$  – integer variable
- Can only be accessed via two indivisible (atomic) operations:
  - **acquire()** and **release()**
    - Originally called **P()** and **V()**
- Less complicated

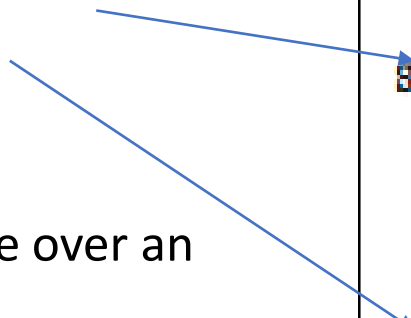
```
acquire() {  
    while value <= 0  
        ; // no-op  
    value--;  
}  
  
release() {  
    value++;  
}
```



# Semaphore Usage

- Associate a semaphore **sem** with a critical section
- To enter a critical section, thread acquires **sem**
- To leave CS, thread releases **sem**
- **Two types:**
  - **Counting** semaphore – integer value can range over an unrestricted domain
  - **Binary** semaphore – integer value can range only between 0 and 1; can be simpler to implement
    - Also known as **mutex locks**
- How do we implement a Semaphore?

```
Semaphore sem = new Semaphore(1);  
  
sem.acquire();  
  
    // critical section  
  
sem.release();  
  
    // remainder section
```



# Java Example Using Semaphores

```
public class Worker implements Runnable {  
    private Semaphore sem;  
    public Worker(Semaphore sem) {  
        this.sem = sem;  
    }  
    public void run() {  
        while (true) {  
            sem.acquire();  
            criticalSection();  
            sem.release();  
            nonCriticalSection();  
        }  
    }  
}
```

```
public class SemaphoreFactory{  
    public static void main(String args[]) {  
        Semaphore sem = new Semaphore(1)  
  
        Thread[] bees = new Thread[5];  
        for (int i = 0; i < 5; i++)  
            bees[i] = new Thread(new Worker(sem));  
  
        for (int i = 0; i < 5; i++)  
            bees[i].start();  
    }  
}
```

# Semaphore Implementation – No Busy waiting

- With each semaphore there is an associated waiting queue. Each entry in a waiting queue has two data items:
  - value (of type integer)
  - pointer to next record in the list
- Two operations:
  - **block** – place the process invoking the operation on the appropriate waiting queue.
  - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue.

# Semaphore Implementation

```
public class Semaphore{  
    private int value;  
    public Semaphore(int value) {  
        this.value = value;  
    }  
    ....  
}
```

```
public synchronized void release() {  
    ++value;  
    notify();  
}
```

- Wake up a waiting process

```
public synchronized void acquire() {  
    while (value <= 0) {  
        try {  
            wait();  
        }  
        catch (InterruptedException e) { }  
    }  
    value--;  
}
```

- Avoid busy waiting by putting the thread to sleep
- System can use a queue and block all threads that are waiting

# Semaphore Implementation

- Must guarantee that no two processes can execute **acquire ()** and **release ()** on the same semaphore at the same time
- Thus, implementation becomes the critical section problem where the wait and signal code are placed in the critical section.
  - Could now have busy waiting in critical section implementation
    - But implementation code is short
    - Little busy waiting if critical section rarely occupied
- Note that applications may spend lots of time in critical sections and therefore this is not a good solution

# Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let **S** and **Q** be two semaphores initialized to 1
  - When p0 executes q.acquire
  - it must wait until p1 executes q.release
  - Similarly, when p1 executes S.acquire
  - It must wait until p0 executes S.release
- **Starvation** – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.

$P_0$	$P_1$
S.acquire();	Q.acquire();
Q.acquire();	S.acquire();
.	.
.	.
.	.
S.release();	Q.release();
Q.release();	S.release();

